

PLATFORM AI TRANSFORMATION PLAYBOOK · ISSUE 01 OF 06

An agent is a production system.

If it can call an API, it sits on the same path as CI, IAM
and on-call.

01

AR

Akhilesh Ranjan Singh

Senior DevSecOps Engineer · github.com/iarsingh

Swipe →

THE WRONG FIRST QUESTION

“Where can we add AI?”

That question produces demos. The better one: which workflow already has an owner, a rollback, and a blast radius we can name?



Chat window vs production system

CHAT WINDOW

- Answers a person
- Failure is a bad answer
- Borrows the user's session
- Ends when the tab closes

AGENT

- Acts for a process
- Failure is a bad write
- Holds its own credentials
- Runs while you sleep

Three gates before any model choice

01 Owner

A named team that gets paged when the agent misbehaves.

02 Rollback

A tested way to undo what the agent changed.

03 Blast radius

Which systems, data, spend and users it can touch, written down.

What it inherits from platform engineering

01 Versioning

Model, prompt and tool set pinned together, like an image tag.

02 Pipeline

Changes ship through review and CI, not a settings page.

03 SLOs

Task success rate, latency, cost per run.

04 On-call

A runbook and a kill switch before the first real user.

Lifecycle I would use

- 1 Lab
- 2 Read-only pilot
- 3 Propose mode
- 4 Owned service
- 5 Review or retire

Each step needs evidence to move forward. Most agents should stay in Propose for a long time.



KEY TAKEAWAY

Treat it like software you would page on.

- Owner, rollback, blast radius first
- Version model + prompt + tools together
- SLOs and a kill switch before users

YOUR TURN

What would you refuse to let an agent touch in your environment?

FOLLOW AKHILESH RANJAN SINGH

Next in the series: #02 GenAI vs Agentic: what actually changes

[Save](#) · [Repost](#)