

PLATFORM AI TRANSFORMATION PLAYBOOK · ISSUE 02 OF 06

GenAI vs Agentic: what actually changes.

The extra boxes in the diagram are software you now have to run.

02

AR

Akhilesh Ranjan Singh

Senior DevSecOps Engineer · github.com/iarsingh

Swipe →

Generative AI

1 Prompt

2 LLM

3 Response

A human reads the output and decides. Failure mode: a wrong answer.
Controls: content filters and answer-quality evals.

Agentic AI

1 **Goal**

2 **Reason**

3 **Plan**

4 **Tools**

5 **Action**

6 **Validate**

Software decides and acts, then loops. Failure mode: a wrong write that other systems treat as fact.

Side by side

	GENAI	AGENTIC
Output	Text	State change
Who acts	A human	Software
Identity	User session	Own workload identity
Testing	Answer evals	Task evals + tool contracts
Operations	Prompt tweaks	SLOs, traces, on-call

The control plane you inherit

01

Identity & authorization

02

Tool & API contracts

03

Data & retrieval quality

04

Memory & state

05

Retries & compensation

06

Traces, logs, audit

07

Human approval

08

Policy & guardrails

THE USEFUL QUESTION

Which process are we willing to redesign?

Because software can now understand, decide and act — with a named owner and a rollback.



KEY TAKEAWAY

Maximum autonomy is not the goal. Appropriate autonomy is.

- An agent is a service in a business process
- Budget for the control plane, not just the model
- Read-only with a gate beats write-capable and unexplainable

YOUR TURN

Harder for your team: building a clever agent, or writing the boundaries it must not cross?

FOLLOW AKHILESH RANJAN SINGH

Next in the series: #03 Identity is the first tool

Save · Repost

