

PLATFORM AI TRANSFORMATION PLAYBOOK · ISSUE 05 OF 06

GitOps, evals and human gates.

An agent should change systems the same way an engineer does.

05

AR

Akhilesh Ranjan Singh

Senior DevSecOps Engineer · github.com/iarsingh

Swipe →

The anti-pattern: chat-to-cluster

01 **kubectl from a prompt**

No diff, no reviewer, no record.

02 **Drift by design**

Cluster state no longer matches Git.

03 **No replay**

Nobody can explain the change at 2am.

The path I want

- 1 Plan as PR
- 2 Policy checks
- 3 Eval gate
- 4 Human approve
- 5 GitOps sync
- 6 Trace

Nothing here is new. It is the path we already use for Terraform and Helm, with an eval step added.

Evals are unit tests for intent

01 Golden tasks

Fixed scenarios with known good outcomes.

02 Tool-call correctness

Right tool, right arguments, right target.

03 Grounding

Claims cite the source they came from.

04 Refusals and budgets

Says no when it should; stays inside cost and latency.

Policy gates worth reusing

01

OPA / Gatekeeper

02

SAST on generated code

03

SCA on new dependencies

04

DAST on changed endpoints

05

Terraform plan checks

06

Secret scanning

07

CODEOWNERS reviewers

08

Signed commits

When a human must approve

01 Irreversible changes

Deletes, migrations, data writes.

02 IAM or network

Anything that widens access.

03 Spend above a threshold

Scaling, new resources.

04 First runs of a new tool

Trust is earned per tool.

KEY TAKEAWAY

If the path to production is a chat message, you skipped the control plane.

- Agents open PRs, GitOps applies
- Evals gate merges like tests do
- Humans approve the irreversible

YOUR TURN

What is your eval for an agent run today — none, a rubric, or a human reading output?

FOLLOW AKHILESH RANJAN SINGH

Next in the series: #06 What I would not put in production

[Save](#) · [Repost](#)