

MLOPS IN PRODUCTION · ISSUE 08 OF 08

Vertex AI vs self-managed GKE.

Pick the platform your team can operate at 2am.

08

AR

Akhilesh Ranjan Singh

Senior DevSecOps Engineer · github.com/iarsingh

Swipe →

The trade-off

	VERTEX AI	GKE
Ops effort	Low	High
Control	Managed defaults	Full
Cost model	Pay per use	You pack the nodes
Custom serving	Within supported runtimes	Anything
Portability	GCP	Any Kubernetes

When I would choose

VERTEX AI

- Small team, no platform squad
- Standard frameworks
- Managed pipelines and endpoints
- Speed over control

GKE

- Existing platform team
- Custom serving or GPU sharing
- Multi-cloud requirements
- Cost at steady high volume

A common hybrid

1 Vertex Pipelines train

2 Model registry

3 Image build

4 GKE serving

5 Shared monitoring

Managed training, self-managed serving where latency and cost need control.

DECISION RULE

Choose by on-call, not by feature list.

A managed service your team understands beats a flexible one nobody can debug.



ML deployment checklist

☒ Model versioned and approved

☒ Eval report vs champion

☒ Canary plan

☒ Rollback by alias

☒ Monitoring signals live

☒ Drift baseline captured

☒ Input data validated

☒ Resource limits set

☒ Owner and on-call named

☒ Model card published

KEY TAKEAWAY

Managed when you can, self-managed when you must.

- Compare ops effort, control and cost
- Hybrids are common and fine
- Use the checklist either way

YOUR TURN

Which did your team choose — and would you choose it again?

FOLLOW AKHILESH RANJAN SINGH

Series complete — save it for your next review

Save · Repost

